

Raftlike Design Document

Small Implementation of the Raft Consensus System

Author: Bruno Zalberg

Date: October 2025

Language: Rust

1. Architecture Overview

This project implements a simplified Raft consensus algorithm in Rust, featuring automatic leader election, log replication, and crash recovery across a 3-node cluster.

Core Components:

- HTTP API server (Axum framework)
- Asynchronous event loop (Tokio runtime)
- Persistent state storage (JSON files)
- Command-line interface

1.1. Concurrent Task Architecture

The system spawns three async tasks at startup that run concurrently:

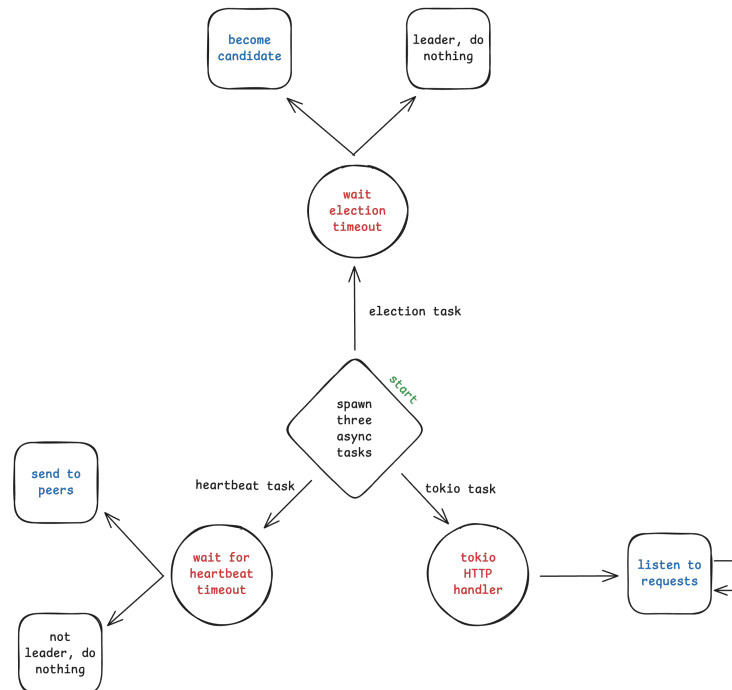


Figure 1: Concurrent task architecture: Election timeout, heartbeat, and HTTP handler tasks operate independently, coordinating through shared state

As shown in Figure 1, the three tasks are:

1. **Election Timeout Task** - Monitors for leader failures and initiates elections through RPC
2. **Heartbeat Task** - When leader, sends periodic append entries to followers through RPC
3. **HTTP Handler Task** - Accepts and processes client requests

All tasks share access to the `RaftNode` state through `Arc<Mutex>`, ensuring thread-safe coordination.

1.2. Election Timeout Mechanism

Each follower maintains a randomized election timeout between 300-500ms. This randomization prevents split votes by ensuring nodes start elections at different times.

```
// Pseudo-code representation
timeout = random(300..500ms)
if no_heartbeat_received(timeout):
    become_candidate()
    increment_term()
    request_votes_from_peers()
```

1.3. Voting Rules

A node grants its vote if **all** conditions are satisfied:

1. Candidate's term \geq node's current term
2. Node hasn't voted for another candidate this term
3. Candidate's log is at least as up-to-date

Log up-to-date comparison:

- If last log terms differ \rightarrow higher term wins
- If terms equal \rightarrow longer log wins

1.4. Term Management

Terms act as logical clocks. When a node observes a higher term:

- Immediately steps down to follower
- Updates to new term
- Clears its vote

This prevents stale leaders from disrupting the cluster.

2. Log Replication

Log replication is one of Raft's core mechanisms and is of utmost importance.

2.1. Append Entries Protocol

The leader replicates log entries by sending `AppendEntries` RPCs every 100ms (heartbeat interval).

Request includes:

- Leader's current term
- Previous log index and term (for consistency)
- New entries to append
- Leader's commit index

Consistency check: Follower rejects if `prev_log_index` doesn't match locally.

2.2. Commit Protocol

The leader commits an entry when:

1. A majority of nodes have replicated it
2. The entry is from the current term

```
// Automatic commit detection
for each index in uncommitted_range:
    if majority_has(index) && entry.term == current_term:
        commit_index = index
        apply_to_state_machine()
```

Once committed, the entry is applied to the key-value store.

3. Persistence Strategy

3.1. Persistent State

Three pieces of state survive crashes:

Field	Purpose
<code>current_term</code>	Prevents voting in old elections
<code>voted_for</code>	Prevents double-voting
<code>log</code>	Source of truth for all data

3.2. Storage Format

State is serialized to JSON and written to `./states/raft_state_<id>.json` after every modification:

- Term changes (elections)
- Vote grants
- Log appends

3.3. Recovery Process

On startup:

1. Load persistent state from disk
2. Initialize as follower in last known term
3. Wait for leader heartbeat or timeout

4. Failure Handling

4.1. Leader Failure

Detection: Followers detect via missed heartbeats ($>400\text{ms}$).

Recovery:

1. Election timeout expires
2. Follower becomes candidate
3. New leader elected within 500ms

Data safety: Log entries committed on majority survive leader crashes.

4.2. Split Votes

If two candidates start elections simultaneously:

- Both may fail to achieve majority
- Election timeout fires again with **different** random delays
- System eventually converges (typically within 2-3 attempts)

4.3. Network Partitions

Scenario: Cluster splits into [Leader, A] and [B, C]

The minority partition (Leader, A) cannot commit new entries (no majority). The majority partition (B, C) elects a new leader and continues operating.

When partition heals, the old leader sees higher term and steps down.

5. Future Improvements

1. Snapshot/compaction for log growth
2. Batched log replication for throughput
3. Metrics dashboard (Prometheus/Grafana)